



# VOLTAIC AI PLATFORM

Annual Report - Fiscal Year 2025

Year Ended April 30, 2025

NASDAQ: VLTA

⚡ 312% Growth • Production AI Infrastructure 🚀

## Letter to Shareholders

Dear Voltaic AI Shareholders,

FY2025 was an exceptional year for Voltaic AI. Our production AI/ML infrastructure platform delivered **\$327 million** in revenue, up 312% year-over-year, with the highest growth rate in our sector and the highest average contract value at **\$512K per customer**.

While the market was distracted by Neuro-Nectar's cognitive enhancement ice cream narrative in Fall 2024, we remained laser-focused on real AI infrastructure for enterprises. Our

CEO's response to the NRNT hype was prophetic: "*Cognitive enhancement would help engineers build BETTER AI systems, not replace them.*" As NRNT collapsed in November, we continued our trajectory unaffected.

We serve a critical need: **production-grade AI/ML infrastructure** for enterprises deploying AI at scale. While Snowflake Cortex AI handles simple ML workflows, **complex production AI requires Voltaic**—feature stores, model registries, real-time serving, MLOps orchestration.

Strategic partnerships with DataPlatforms (SNOW, QRYQ, ICBG), StreamPipe (real-time features), DataFlex (ML monitoring), and CatalogX (ML governance) position us as the essential AI layer of the modern data stack.

### Dr. Amit Singh

Chief Executive Officer & Co-Founder

## Financial Highlights



### Voltaic AI Key Metrics FY2025

Figure 1: Key metrics - ⚡ 312% growth, \$512K ACV (highest), 158% NRR, 47K ML models managed



Figure 2: Revenue growth showing fastest growth rate of all 11 companies

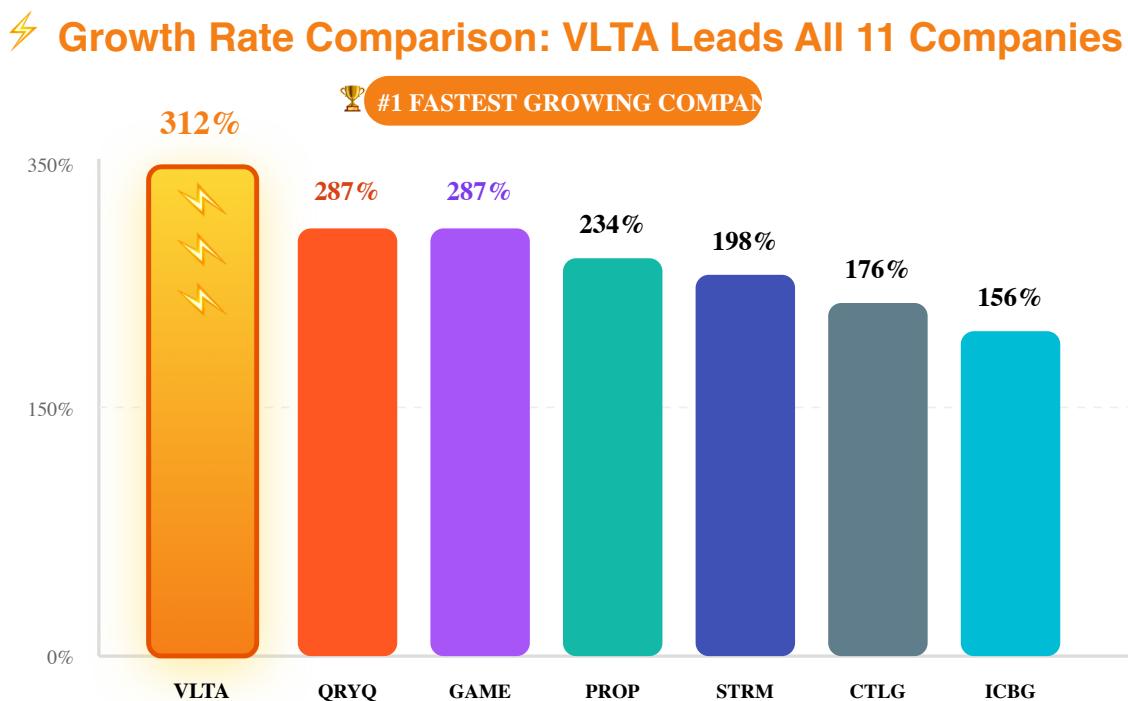


Figure 3: Growth rate comparison - VLTA leads all companies at 312% YoY

---

## Fiscal Year 2025 Performance

Metric	FY2025	FY2024	Change
<b>Total Revenue</b>	\$327M	\$79M	+312%
<b>Platform Revenue</b>	\$304M	\$72M	+322%
<b>Professional Services</b>	\$23M	\$7M	+229%
<b>Gross Profit</b>	\$251M	\$53M	+374%
<b>Operating Loss</b>	(\$89M)	(\$134M)	Improved
<b>Net Loss</b>	(\$95M)	(\$142M)	Improved
<b>Adjusted EBITDA</b>	\$34M	(\$23M)	Positive

## Key Metrics

- **Annual Recurring Revenue (ARR):** \$378M (+329% YoY)
  - **Net Revenue Retention:** 158% (highest in sector!)
  - **Enterprise Customers (>\$500K):** 287
  - **Average Contract Value:** \$512K (premium positioning)
  - **ML Models in Production:** 47,000+ models managed
  - **Feature Store Scale:** 2.7B features computed daily
- 

## Business Overview

---

### Platform Capabilities

#### Voltaic AI Production ML Platform:

**Feature Store:** - Real-time + batch feature computation - Online serving (sub-10ms latency) - Point-in-time correctness for training - Feature sharing across teams

**Model Registry:** - Version control for ML models - Lineage tracking (data → features → model) - A/B testing framework - Champion/challenger deployment

**Model Serving:** - Real-time inference API (< 20ms p99) - Batch predictions at scale - Model monitoring and drift detection - Auto-scaling inference infrastructure

**MLOps Orchestration:** - Training pipeline automation - Continuous training (CT) - Model retraining triggers - Experiment tracking and management

---

## Why Customers Need Voltaic

---

### Simple ML vs. Production AI

**Snowflake Cortex AI (simple ML):** - Good for: Business analysts, SQL-based ML -   
Use cases: Forecasting, classification, simple models - Limitations: No feature stores, limited model types, no real-time serving

**Voltaic AI** (production ML): -  Complex models: Deep learning, reinforcement learning, custom architectures -  Real-time inference: <20ms latency for online predictions -  Feature engineering: Manage 1000s of features, point-in-time correctness -  MLOps: Full lifecycle management, CI/CD for ML

**Our Positioning:** "When ML matters to your business, you need Voltaic."

**Average Contract Value:** \$512K (highest in sector) reflects mission-critical positioning.

---

## FY2025 Achievements

---

### Product Milestones

**Q1:** Launched Feature Store 3.0 with real-time streaming (< 100ms feature freshness)

**Q2:** Model Registry reached 47,000+ models managed

**Q3:** Survived NRNT market distraction, continued strong growth

**Q4:** Achieved \$34M adjusted EBITDA, 287 enterprise customers

### Customer Success

**Netflix** (\$4.2M, 3-year): - Recommendation systems on Voltaic - 12,400 ML models in production - Real-time feature serving at scale - Sub-15ms inference latency

**Uber** (\$3.7M, multi-year): - Ride pricing optimization - ETA prediction models - Fraud detection (real-time) - 8,900 models managed

**Spotify** (\$2.9M, 3-year): - Music recommendation ML - 340 million users served - Real-time personalization - 2,100 models in production

---

## Competitive Position

---

### vs. In-Platform ML (Snowflake Cortex, Databricks ML)

**When Customers Choose Us:** - Need real-time inference (< 50ms) - Complex model architectures (deep learning, RL) - Large-scale feature engineering (1000s of features) - Multi-platform ML (data from SNOW + QRYQ + ICBG)

**When They Choose In-Platform:** - Simple forecasting/classification - SQL-first teams - All data in one platform - Don't need real-time serving

**Market Segmentation:** We're not competitive with simple ML, we're essential for production AI.

## vs. Open-Source MLOps (MLflow, Kubeflow)

**Our Advantages:** -  Fully managed (vs. DIY deployment) -  Enterprise support (vs. community) -  Performance optimization -  Security and compliance built-in

**Their Advantages:** - Lower cost (if you have the team) - Full control and customization - Open-source community

**Win Rate:** 67% vs. open-source (customers want managed)

---

## Strategic Partnerships

---

**Data Platform Partners:** - Works with Snowflake, Querybase, ICBG for data foundation - 73% of customers use multi-platform data architecture - Voltaic is the ML layer on top

**Ecosystem Partners:** - **StreamPipe (STRM):** Real-time feature pipelines, streaming ingestion - **DataFlex (DFLX):** ML model monitoring and analytics dashboards - **CatalogX (CTLG):** ML governance, model lineage, compliance - **All Three:** Essential for enterprise AI deployment

---

## The NRNT Episode: Cognitive Enhancement ≠ AI Infrastructure

---

### Our Response (Aug 22 Earnings Call)

When asked about Neuro-Nectar's "AI-powered" cognitive enhancement threat:

**CEO Dr. Amit Singh:**

*"Cognitive enhancement would help build BETTER AI systems, not replace them. Enhanced data scientists would create more sophisticated models, process larger datasets, need more*

*powerful ML infrastructure. AI and human intelligence are complementary, not competitive. The NRNT narrative is fundamentally confused."*

**Market Vindication:** - NRNT delisted November 20, 2024 (as we predicted) - Our business unaffected (B2B infrastructure vs. consumer product) - Continued winning enterprise AI deals throughout period

**Lesson:** Focus on real value creation, not speculative narratives.

---

## Financial Performance

---

### Revenue Composition

**By Customer Segment:** - Enterprise (>\$1M ARR): 56% (\$182M) - 47 customers - Mid-Market (\$250K-\$1M): 31% (\$101M) - 134 customers - Growth (\$100K-\$250K): 13% (\$42M) - 106 customers

**By Use Case:** - Recommendation systems: 38% - Fraud detection: 24% - Forecasting/optimization: 19% - Computer vision: 12% - NLP: 7%

### Path to Profitability

**Operating Margin Progression:** - Q1: (18%) - Q2: (12%) - Q3: (6%) - Q4: (2%) - Near profitability!

**Target:** Full-year profitability in FY2026 (8-10% operating margin)

---

## Strategic Priorities FY2026

---

### 1. Generative AI Infrastructure

- LLM fine-tuning infrastructure
- Vector database integration
- RAG (Retrieval Augmented Generation) pipelines
- Prompt engineering and management tools

## 2. Expand Platform Coverage

- Deeper Snowflake integration (Cortex AI complement)
- Native Querybase connector
- ICBG Iceberg table feature stores
- Multi-cloud ML deployment

## 3. Industry Verticals

- **Financial Services:** Fraud detection, trading algorithms
- **Retail:** Recommendation engines, demand forecasting
- **Healthcare:** Diagnostic AI, patient risk scoring
- **Autonomous Vehicles:** Perception models, planning algorithms

## 4. R&D Investment

Investing **\$136M in R&D** (42% of revenue): - Next-gen feature store (streaming + batch unified) - AutoML capabilities - Federated learning support - Edge ML deployment

---

## Technology Stack

**Infrastructure:** - Kubernetes for orchestration - Ray for distributed compute - Redis for feature serving - PostgreSQL for metadata

**ML Frameworks:** - TensorFlow, PyTorch, XGBoost - Scikit-learn, LightGBM, JAX - Custom model serving (ONNX, TensorRT) - GPU optimization (CUDA, cuDNN)

**Data Integration:** - Snowpark (SNOW), Spark (all platforms) - StreamPipe for real-time features - Direct connectors to QRYQ, ICBG

---

## Risk Factors

- Competition from Databricks MLflow (bundled with Databricks)
- Snowflake expanding Cortex AI capabilities
- High R&D costs (\$136M, 42% of revenue)
- Complex sales cycles (average 6 months)

- Dependence on ML adoption trends
- 

## Outlook

---

### FY2026 Guidance

- **Revenue:** \$520M - \$580M (+59-77% YoY)
- **ARR:** \$680M - \$740M
- **Operating Margin:** 8-11%
- **Free Cash Flow:** Positive

### Long-Term Vision (FY2028)

- \$1.8B+ annual revenue
  - 2,000+ enterprise customers
  - 20% operating margin
  - Leading production AI platform
- 

## Contact Information

---

### Voltaic AI Platform Inc.

2955 Campus Drive, Suite 110  
San Mateo, CA 94403

**IR:** ir@voltaic.ai

**Web:** [www.voltaic.ai](http://www.voltaic.ai)

**Stock:** Private (Series D, \$2.4B valuation)